

YOURSTORY



II2S



Coco CLI

HACKATHON

GCC Edition



Customer 360 Decisioning Platform

Team Name : Simplifiers

Team Leader Name : Sairaj Sawant

Team Size : 1

Problem Statement : Customer 360 and Next Best Action Engine

Customers signal what they'll do long before they do it

Insurers and lenders already hold the evidence. It's spread across systems nobody reads together.



Scattered data

Policies, claims, payments, calls, emails and tickets live in separate systems.



Missed warnings

A customer threatens to switch insurers on a call while their claim sits stuck for weeks. Nobody links the two.



Inconsistent action

Every RM decides alone, and nobody learns which actions actually work.

WHO IT'S FOR



Relationship Manager

A ranked worklist and the next best action for each customer



Team Lead

Approves larger offers for the team



Analyst / Domain Expert

Tunes signals, weights and rules without code

Indian health insurance and retail lending: IRDAI grievances, cashless claims, portability, EMI stress

From customer question to next best action

Unify every touchpoint, read it with AI, decide with approved rules, and learn from every outcome.

1

Unify

Calls, emails, tickets, payments and policies become one live customer view.

510

customers unified, every one scored

8

source types, in English and Hinglish

2

Understand

AI picks out intent, sentiment and evidence from every conversation.

2,796

live signals, each with its evidence

3

decision engines running side by side

3

Decide

Approved rules rank the next best action for each role.

4

reusable CoCo CLI skills

19

config rows to add a new engine


4

Learn

Every outcome updates which actions work best.

Architecture: four independent layers

Each layer reads only from the one before it, so any layer can change without touching the rest.



Layer 1 · Every touchpoint, one customer

510 Indian customers. Dynamic Tables keep it current; a semantic view gives everyone the same definitions.

SOURCES · RECORDS

CONVERSATIONS

Call transcripts	295
Emails	2,058
Support tickets	1,523

TRANSACTIONS

Policies & versions	4,269
Payments	1,903
Claims	215
Loans	15

REFERENCE

Product documents	11
Regulatory filings	92

UNIFY

Canonical model

Dynamic Tables keep it current

Customer

Account

Product

Interaction

Event

SERVE

Customer 360 profile

Relationship value, tenure, segment
Open tickets, SLA breaches, CSAT
Claims, grievances, portability
Every call and email, in context
One dated timeline

SEMANTIC VIEW

Governed definitions for analytics and the agent

CHANGE STREAMS

New conversations trigger signal extraction


+ New source: map it once to the model

Signals, decisions and apps keep running untouched

510 customers in English and Hinglish: CoCo CLI seeded 30 from Kaggle and IRDAI samples, then scaled to 510 from those

Layer 2 · AI turns conversations into evidence

AI extracts facts into a fixed vocabulary, and every signal keeps the sentence that justifies it.



25

signals live, 2,796 readings,
each with its evidence

SIGNAL DISCOVERY · RUNS DAILY

Scan unused data


Propose signals

Expert approves

Live everywhere


Layer 3 · Isolated engines on a shared vocabulary

Each use case is its own engine. All read the same signals, and none can change another.




Repeatable recommendations that keep improving

AI reads, rules decide. Scoring runs on measured track records, so every answer can be reproduced and explained.



Proactive by default, with people in the loop

The platform watches every customer on a schedule and routes each decision to the person with the right authority.



Adding a use case takes configuration only

We added Service Recovery to the working platform. Nothing that already existed changed.

19

configuration rows

0 lines of engine code

Existing engines returned identical results before and after the change.

BEFORE


- 3 customers with stuck claims were on nobody's worklist
- 2 customers we had let down were offered a loan top-up

AFTER

- All 3 now top the list with an expedited claim
- Both get a priority callback, and no upsell until it's fixed

Built and run with CoCo CLI

CoCo CLI seeded the data, built the platform, and publishes four reusable skills to Snowflake for any team.



Snowflake capabilities in use

Data, AI and the app all run inside Snowflake. No customer data leaves the account.



Dynamic Tables

Keep the unified customer view current



Streams & Tasks

Detect change and run the pipeline on schedule



Snowpark Python

Portfolio-wide scoring and signal discovery



Semantic View

Shared, governed business definitions



AI_COMPLETE

Intent from calls, held to a fixed list



AI_SENTIMENT

Tone of every conversation



AI_SUMMARIZE

A customer's history in one paragraph



Cortex Search

Search calls, emails and product documents



Cortex Agent

Natural-language access to every tool



Streamlit in Snowflake

The app runs right next to the data



Masking & row access

PII masked, each persona sees only its book



CoCo CLI

Built the platform and packages it as skills

What's live in the app today

Two modes on live Snowflake data: a guided studio and the everyday console.

Scenario Studio

The full loop on any customer, step by step

- 1 **Stage an event** describe it, AI writes a realistic call
- 2 **Detect** signals extracted with their evidence
- 3 **Understand** state recomputed, conflicts resolved
- 4 **Decide** ranked for your role, limits checked
- 5 **Act** carried out, customer notified, call brief
- 6 **Learn** outcome recorded, ranking re-run

Operations Console

Each role sees only the pages it's allowed to use

- **My Feed** today's ranked workload
- **Decision Queue** by severity and value
- **Approvals** one-click manager sign-off
- **Customer 360** next best action, act in place
- **Portfolio & Learning** value at risk, what works
- **Ask the data** chat that remembers context
- **Config Studio** weights, limits, rules
- **Signal Discovery** daily suggestions

Enterprise ready by design



Private by default

Data stays in your account, and masking policies hide PII from roles that don't need it.



Scoped by role

Row access policies: RMs see their own book, team leads their team, analysts all.



Spending limits

Offers above a role's limit wait for someone with the authority.



Signals with evidence

Each signal keeps the quote and confidence behind it.



Repeatable decisions

Same input, same answer. AI reads, rules decide.



Full audit trail

Every recommendation, action and outcome is recorded and reversible.

Configuration-driven. Weights, limits, rules, signals and whole use cases change without a release.

Impact

Measured on the prototype's live data.

1,245

of 2,796 signals come only from what customers said, invisible to structured systems

38

at-risk customers caught only from what they said, with nothing in their records

~2 sec


to rank all 510 customers into today's worklist

1 question

from a customer question to a recommended action


In production: less prep per customer, consistent actions across RMs, and a measured view of which actions work.

Where the 2,796 signals come from



New use cases in days, new domains in weeks

Everything above the data mapping is reused. A domain expert shapes the ontology; the platform does the rest.



MUTUAL FUND CUSTOMER 360, SAME PLATFORM

SIP stop risk

Redemption intent on calls

Tax-season fund fit

Failed mandate recovery

An incremental 360 that scales at marginal cost

Work is done only where something changed, so a bigger book doesn't mean a bigger bill.



Incremental Dynamic Tables

Facts refresh only where the data changed



Changed-customer streams

Signals and state recompute only for customers who changed



Daily time tick

Date-based signals touch only customers crossing a threshold today



Signal fingerprints

Customers whose signals didn't change are skipped

AI cost grows with new conversations. Rule cost grows with changed customers, and time-based signals with threshold crossings.

Beyond the demo

Each step builds on the platform as it is. Estimates assume one engineer.



Connect real sources

CRM, core policy and loan systems and the contact centre through Snowflake connectors, mapped once to the canonical model.

1–2 days per source



Actions in Slack and Jira

Approved actions post to the RM's channel and open a ticket automatically.

2–3 days



Fine-grained access

Extend row access by region and branch, and masking to every new source.

2–3 days



Cross-domain, any entity

The same engines over households, SME accounts or policies, with shared suppression across all of them.

1–2 weeks



AI onboarding wizard

Describe a use case or signal in plain English. AI drafts the configuration and a person approves it.

1–2 weeks



Predictive ML signals

Snowflake ML forecasts which customers are trending toward decline and which product they'll need next, before it shows in a call.

1–2 weeks

What the demo shows

One workflow end to end, input to output, in CoCo CLI and then in the app.

INPUT

An RM asks CoCo CLI

“What should we do about Suresh Reddy?” and then a follow-up about another customer.



PROCESSING

Customer Query skill

Finds the customer, picks the one engine that answers the question, and runs it on Snowflake.



OUTPUT

A grounded answer

The ranked action with its track record and approval status, backed by his own calls.

THEN IN THE APP

My Feed across all three engines · act straight from Customer 360 · a chat that remembers context · the full loop in Scenario Studio · one-click manager approval

YOURSTORY



1125

Coco CLI

HACKATHON

GCC Edition

THANK YOU

